

IVrixDB Tests Overview

KEY THEME: Each test builds upon all the aspects and features of the previous tests. This means that the tests gradually become more **integrated** and ensures that the system can **function at scale**.

1. Indexing into IVrix Cluster
 - a. Direct indexing into IVrix Node
 - b. Load-Balanced Indexing into IVrix Cluster
 - c. Node-Failure during Load-Balanced Indexing
 - i. Forwarder Failure
 - ii. Forwarder Failure
 - d. Multiple batches for same index on one IVrix Node
 - e. segregation of data into buckets with proper index rollover
2. Single Search on IVrix Cluster
 - a. Hold Bucket, then Release once finished with Bucket
 - b. Time-Ranged Search
 - i. Hot
 - ii. Warm
 - iii. Cold
 1. attach/detach of COLD buckets
 - c. Unbounded Search (Search Features in General)
 - d. Keyword Search
 - e. Statistics Search (Results page)
 - f. Sort Search
 - g. Sort + Statistics Search
 - h. "SELECT" solr-function Search
 - i. "HAVING" solr-function Search (filtering on field from STFE)
 - j. Chained Search (w/ streaming and non-streaming commands)
3. Multiple Searches on IVrix Cluster
 - a. Searches waiting on one another
4. Indexing + Single Search on IVrix Cluster (index rollover during search)
 - a. All Cold being held during index rollover
 - b. Oldest Warm being held during index rollover
5. Indexing + Multiple Searches on IVrix Cluster
6. Node Failure during Indexing + Searching on IVrix Cluster
 - a. Node Failure
 - b. Overseer-specific Failure
 - c. Slave-specific Failure
 - d. OPTIONAL TESTS --- operation failure/recovery/partial-failure/completion
7. Recovery of Node after Node failure during Indexing + Searching on IVrix Cluster

Tests Details

Setup for All Tests (unless explicitly stated for each point)

- Two IVrix Nodes, One Zookeeper Node
- 8.03 million mock logs, spanning over 11 years with 2000 events per day
 - o As one file
 - o As two files
 - Split into two files, in a round-robin batched separation method
- Config in "Constants" class:
 - o DEFAULT_REPLICATION_FACTOR of 2 (for buckets)
 - o MAXIMUM_ATTACHED_HOT_PLUS_WARM_BUCKETS_PER_NODE = 10
 - o MAXIMUM_ATTACHED_COLD_BUCKETS_PER_NODE = 3
 - o DEFAULT_MAX_EVENT_COUNT_IN_BUCKET = 401500
 - o DEFAULT_EVENT_BATCH_SIZE = 30,000

Teardown

- Restart IVrix Nodes (not always though). This is done for two reasons that are further detailed in the "Disclaimers" section of "Disclaimers, Warnings, and Bugs" document:
 - o Due to the bug where an IVrix Index cannot be fully deleted without a restart of the residing nodes, the IVrix Nodes must be restarted. Additionally, the directories of detached buckets still exist, and so manual removal of them is a must.
 - o Due to no memory resource manager for search, memory just keeps growing and growing during run-time. That starts to affect the machine itself, and so periodic restarts are necessary.

1 – Indexing into IVrix Cluster

1A – Direct indexing into IVrix Node

- **TEST 1:**
 - o Procedures:
 - Index the one file with 8.03 million logs into one IVrix Node (w/o load-balancing)
 - o Expected outcome:
 - One Bucket with 8.03 million logs

1B – Load-Balanced Indexing into IVrix Cluster

- **TEST 1:**
 - o Procedures:
 - Index the one file with 8.03 million logs into one IVrix Node (w/ load-balancing)
 - o Expected outcome:
 - 20 buckets, 10 per node
 - 420k events per bucket (2_0k on last two buckets). Surpasses limit because of batch sizing

1C – Node-Failure during Load-Balanced Indexing

Forwarder Failure

- **TEST 1:**
 - Procedures:
 - Index the one file with 8.03 million logs into IVrix Node 8983 (w/ load-balancing)
 - Kill IVrix Node 8983 unexpectedly
 - Expected outcome:
 - Forwarder will die. No more indexing from that request.

Forwarder Failure

- **TEST 1 (During Indexing into Dead Node):**
 - Procedures:
 - Index the one file with 8.03 million logs into IVrix Node 8983 (w/ load-balancing)
 - Kill IVrix Node 7574 while forwarder sent batch to it
 - Add breakpoint at EventForwarder.tryToIndexBatch() on .request() line to assist in timing
 - Expected outcome:
 - Forwarder will re-send batch to a live node
 - All Events will be indexed
 - Duplicate Events will likely exist (maximum of 30k duplicates)
- **TEST 2 (During Indexing into Other Node):**
 - Procedures:
 - Index the one file with 8.03 million logs into IVrix Node 8983 (w/ load-balancing)
 - Kill IVrix Node 7574 while forwarder sent batch to 8983
 - Add breakpoint at EventForwarder.tryToIndexBatch() on .request() line to assist in timing
 - Expected outcome:
 - Forwarder will ignore dead node
 - All Events will be indexed

1D – Multiple batches for same index on one IVrix Node

- **TEST 1:**
 - Procedures:
 - Index one half file into IVrix Node 8983 (w/ load-balancing)
 - Index the other half file into IVrix Node 7574 (w/ load-balancing)
 - Expected outcome:
 - 20 buckets, 10 per node
 - 420k events per bucket (2_0k on last two buckets). Surpasses limit because of batch sizing

1E – Segregation of Data into Buckets with Proper Index Rollover

- **TEST 1:**
 - Setup:
 - MAXIMUM_ATTACHED_HOT_PLUS_WARM_BUCKETS_PER_NODE = 10

- MAXIMUM_ATTACHED_COLD_BUCKETS_PER_NODE = 3
 - Procedures:
 - Index the one file with 8.03 million logs into IVrix Node 8983 (w/ load-balancing)
 - Expected outcome:
 - 20 buckets
 - (HOT, WARM, COLD, DETACHED)
 - Node 8983 -> (1, 4, 3, 2)
 - Node 7574 -> (1, 4, 3, 2)
- **TEST 2:**
 - Setup:
 - MAXIMUM_ATTACHED_HOT_PLUS_WARM_BUCKETS_PER_NODE = 4
 - MAXIMUM_ATTACHED_COLD_BUCKETS_PER_NODE = 1
 - Procedures:
 - Index the one file with 8.03 million logs into IVrix Node 8983 (w/ load-balancing)
 - Expected outcome:
 - 20 buckets
 - (HOT, WARM, COLD, DETACHED)
 - Node 8983 -> (1, 1, 1, 7)
 - Node 7574 -> (1, 1, 1, 7)

2 – Single Search on IVrix Cluster

SETUP FOR ALL TESTS IN THIS SECTION:

- Have the index ready and fully indexed
- index the one file with 8.03 million logs into IVrix Node 8983 (w/ load-balancing) [**1B TEST 1, PROCEDURE**]

2A – Hold Bucket, then Release once finished with Bucket

- **TEST:**
 - Procedures:
 - Run un-bounded search
 - Expected outcome:
 - For Every Hold of a group of buckets, there should be a release of them **BEFORE** the hold of the next group of buckets.

2B – Time-Ranged Search

- **TEST 1:**
 - Procedures:
 - Run ranged search over HOT buckets
 - (_default20)
 - latest: 2010-12-29T02:46:40Z, earliest: 2010-12-22T01:23:25Z
 - Run ranged search over WARM buckets

- (_default18, _default17, _default16, _default15)
 - latest: 2009-11-16T02:05Z, earliest: 2008-10-12T00:00:05Z
 - Run ranged search over COLD buckets
 - (_default4, _default3, _default2, _default1)
 - latest: 2002-03-13T01:23:20Z, earliest: 2000-01-02T00:00:05Z
 - Expected outcome:
 - HOT buckets – 15,000 events
 - WARM buckets – 801,500 events
 - COLD buckets – 1,603,000 events
 - Attaching/Detaching buckets appropriately
 - Maintaining attachment constraint limits
- **TEST 2:**
- Procedures:
 - Run ranged search over 1 Timeline bucket
 - latest: 2010-12-01T00:00Z, earliest: 2010-11-01T00:00Z
 - Run ranged search over 2 Timeline buckets
 - latest: 2010-12-01T00:00Z, earliest: 2010-10-01T00:00Z
 - Run ranged search over 4 Timeline buckets
 - latest: 2010-12-01T00:00Z, earliest: 2010-08-01T00:00Z
 - Expected outcome:
 - 1 Timeline bucket – 60,000 events
 - 2 Timeline bucket – 122,000 events
 - 4 Timeline bucket – 244,000 events

2C – Unbounded Search (Search Features in General)

- **TEST:**
- Procedures:
 - Run un-bounded search:
 - search(_default,zkHost="localhost:9983",ivrix=true,qt="/export",q="*:*","fl="_raw,_time,id",sort="_time desc",partitionKeys=_time)
 - Expected outcome:
 - Time range – over a span of 11 years (2000-01-01T00:00Z to 2011-01-01T00:00Z)
 - Number of events – 8,030,000
 - Should attach/detach over COLD buckets when necessary
 - **ALL SEARCH API COMPONENTS ARE INCREMENTALLY UPDATED OVER THE SEARCH PERIOD**
 - Timeline
 - Timeline will grow as events are digested in reverse time order (latest to earliest)
 - At the end
 - Timeline buckets in the scope of “MONTH”

- each bucket has between 56k to 62k events
- Events
 - All-inclusive bucket has only 1k events stored in it
 - Each subsequent bucket has only 1k events stored in it
 - Un-bounded pagination of events returns only 1k events
 - Bounded pagination of events returns **1k*N** events, depending on the number of buckets selected
- Field Summaries
 - Number fields show stats (min, max, avg)
 - String and Number fields show top 10 values (count/percentage of each value)
- Results
 - Empty

2D – Keyword Search

- **TEST:**
 - Procedures:
 - Run un-bounded, keyword search (keyword “simpson”):
 - `search(_default,zkHost="localhost:9983",ivrix=true,qt="/export",q="_raw:simpson",fl="_raw,_time,id",sort="_time desc",partitionKeys=_time)`
 - Expected outcome:
 - Timeline, Events, and Field Summaries just as **2C**, **EXCEPT**:
 - Timeline
 - Number of events – 1,604,004
 - Each bucket has between 10k to 13k events
 - All events must have keyword “simpson” in them

2E – Statistics Search (Results page)

- **TEST:**
 - Procedures:
 - Run un-bounded, statistics search (group by over “age”, sum over “random”):
 - `rollup(search(_default,zkHost="localhost:9983",ivrix=true,qt="/export",q="*:.*",fl="_raw,_time,id",sort="_time desc",partitionKeys=_time),over=age,sum(random))`
 - Expected outcome:
 - Timeline, Events, and Field Summaries just as **2C**
 - Results
 - **INCREMENTALLY UPDATED OVER THE SEARCH PERIOD**
 - Groups over age of 10, 20, 30, 40, 50, 60, and 70
 - Each group has an ever-increasing value of “sum(random)”

2F – Sort Search

- **TEST:**
 - Procedures:
 - Run un-bounded, sort search (top 10k, random desc):

- `top(n=10000,search(_default,zkHost="localhost:9983",ivrix=true,qt="/export",q="*:*",fl="_raw,_time,id",sort="_time desc",partitionKeys=_time),sort="random desc")`
- Expected outcome:
 - Timeline, Events, and Field Summaries just as **2C**, **EXCEPT**:
 - Timeline
 - Number of events – 10,000. This number does **NOT** change as the search goes on
 - Number of events per timeline bucket – changes frequently, and gets smaller and smaller as the search goes on
 - Events
 - Same “page” has a strong likelihood of changing after each refresh

2G – Sort + Statistics Search

- **TEST:**
 - Procedures:
 - Run un-bounded, sort + statistics search (top 10k, sorted by random DESC, THEN sum over random):
 - `rollup(top(n=10000,search(_default,zkHost="localhost:9983",ivrix=true,qt="/export",q="*:*",fl="_raw,_time,id",sort="_time desc",partitionKeys=_time),sort="random desc"),over=age,sum(random))`
 - Expected outcome:
 - Timeline, Events, and Field Summaries just as **2F**
 - Results just as **2E**, except that “sum(random)” is not ever-increasing by large changes, rather it is ever-changing

2H – “SELECT” solr-function Search

- **TEST:**
 - Procedures:
 - Run unbounded, “SELECT” solr-function Search (replace all instances where age=70 with value 1000):
 - `select(search(_default,zkHost="localhost:9983",ivrix=true,qt="/export",q="*:*",fl="_raw,_time,id",sort="_time desc",partitionKeys=_time),age,_time,_raw,weight,random,replace(age,70,withValue=1000))`
 - Expected outcome:
 - Timeline, Events, and Field Summaries just as **2C**, **EXCEPT**:
 - Events where age=70, the value is replaced with 1000

2I – “HAVING” solr-function Search (filtering on field from STFE)

- **TEST:**
 - Procedures:

- Run unbounded, “HAVING” solr-function Search (filtering on STFE field “age”, where age = 10)
 - `having(search(_default,zkHost="localhost:9983",ivrix=true,qt="/export",q="*:*",fl="_raw,_time,id",sort="_time desc",partitionKeys=_time),eq(age,10))`
- Expected outcome:
 - Timeline, Events, and Field Summaries just as **2C**, **EXCEPT**:
 - Timeline
 - Number of events – 1,146,982
 - Each bucket has less than 10k (between 8k to 10k) events
 - Events
 - There are **ONLY** events with age = 10

2J – Chained Search (w/ streaming and non-streaming commands)

- **TEST:**
 - Procedures:
 - Run unbounded, keyword + SELECT + sort + HAVING + statistics search
 - `rollup(having(top(n=10000,select(search(_default,zkHost="localhost:9983",ivrix=true,qt="/export",q="_raw:simpson",fl="_raw,_time,id",sort="_time desc",partitionKeys=_time),age,_time,_raw,weight,random,replace(age,70,withValue=1000)),sort="random desc"),eq(age,10)),over=age,sum(random))`
 - Expected outcome:
 - Timeline
 - Number of events – 1,445
 - Each bucket has less than 20 events
 - Events
 - Have only fields age, _time, _raw, weight, and random
 - There are **ONLY** events with age = 10
 - Results
 - One group (age of 10)
 - “sum(random)” is not ever-increasing by large changes, rather it is ever-changing

3 – Multiple Searches on IVrix Cluster

SETUP FOR ALL TESTS IN THIS SECTION:

- Have the index ready and fully indexed

- index the one file with 8.03 million logs into IVrix Node 8983 (w/ load-balancing) [**1B TEST 1, PROCEDURE**]

3A – Searches waiting on one another

- **TEST:**
 - Procedures:
 - Run 4 ranged searches simultaneously spanning over the COLD buckets (half running on Node 8983, half running on Node 7574)
 - 7574 #1 -- latest: 2006-05-22T02:46:40Z, earliest: 2004-05-23T02:46:40Z
 - 8983 #1 -- latest: 2004-05-22T02:46:40Z, earliest: 2002-05-02T01:23:25Z
 - 8983 #2 -- latest: 2003-04-18T00:41:40Z, earliest: 2001-03-27T02:05:05Z
 - 7574 #2 -- latest: 2002-03-13T01:23:20Z, earliest: 2000-01-02T00:00:05Z
 - Expected outcome:
 - Some Searches will need to wait to attach, since not all buckets can be attached at once
 - 7574 #1 -- 1458001 events
 - 8983 #1 -- 1503000 events
 - 8983 #2 -- 1503000 events
 - 7574 #2 -- 1603000 events

4 – Indexing + Single Search on IVrix Cluster (index rollover during search)

SETUP FOR ALL TESTS IN THIS SECTION:

- MAXIMUM_ATTACHED_HOT_PLUS_WARM_BUCKETS_PER_NODE = 4
- MAXIMUM_ATTACHED_COLD_BUCKETS_PER_NODE = 1

4A – All Cold Being Held During Index Rollover

- **TEST:**
 - Procedures:
 - Add breakpoint at IVrixOverseer.executeCreateBucketOperation() before the state lock command
 - With condition: state.getIndex(indexName).getBucket("_default7") == null && state.getIndex(indexName).getBucket("_default6") != null

- index the one file with 8.03 million logs into IVrix Node 8983 (w/ load-balancing) **[1B TEST 1, PROCEDURE]**
- Once breakpoint is hit, execute search on (_default1, _default2)
 - latest: 2000-07-20T00:00:05Z, earliest: 2000-01-02T00:00:05Z
- Release breakpoint
- Expected outcome:
 - Before Search finishes:
 - _default3 is rolled to cold, and detached
 - _default7 is created

4B – Oldest Warm Being Held During Index Rollover

- **TEST:**
 - Procedures:
 - Add breakpoint at IVrixOverseer.executeCreateBucketOperation() before the state lock command
 - With condition: state.getIndex(indexName).getBucket("_default5") == null && state.getIndex(indexName).getBucket("_default4") != null
 - index the one file with 8.03 million logs into IVrix Node 8983 (w/ load-balancing) **[1B TEST 1, PROCEDURE]**
 - Once breakpoint is hit, execute search on (_default1, _default2)
 - latest: 2000-07-20T00:00:05Z, earliest: 2000-01-02T00:00:05Z
 - Release breakpoint
 - Expected outcome:
 - Waits for search on _default1 to finish
 - In the meantime, no new searches on _default1 are created, even if requested
 - Once search on _default1 finishes
 - _default 1 rolls to cold
 - _default5 is created

5 – Indexing + Multiple Searches on IVrix Cluster

- **TEST:**
 - Procedures:
 - index the one file with 8.03 million logs into IVrix Node 8983 (w/ load-balancing) **[1B TEST 1, PROCEDURE]**
 - Wait until bucket _default16 is made
 - Run 4 ranged searches simultaneously spanning over the COLD buckets (half running on Node 8983, half running on Node 7574)
 - 7574 #1 -- latest: 2006-05-22T02:46:40Z, earliest: 2004-05-23T02:46:40Z

- 8983 #1 -- latest: 2004-05-22T02:46:40Z, earliest: 2002-05-02T01:23:25Z
 - 8983 #2 -- latest: 2003-04-18T00:41:40Z, earliest: 2001-03-27T02:05:05Z
 - 7574 #2 -- latest: 2002-03-13T01:23:20Z, earliest: 2000-01-02T00:00:05Z
- Expected outcome:
 - Some Searches will need to wait to attach, since not all buckets can be attached at once
 - When Indexing needs to roll oldest warm bucket, but it is being held
 - Waits for search on bucket to finish
 - In the meantime, no new searches on bucket are created, even if requested
 - When Indexing needs to make space for cold bucket, but all cold is being held
 - Detaches that cold bucket
 -
 - 7574 #1 -- 1458001 events
 - 8983 #1 -- 1503000 events
 - 8983 #2 -- 1503000 events
 - 7574 #2 -- 1603000 events
 - Indexing will finish with all 20 buckets filled, and all 8.03 million logs

6 – Node Failure during Indexing + Searching on IVrix Cluster

- **TESTS 1 & 2:**
 - Procedures:
 - index the one file with 8.03 million logs into either Node 8983 or Node 7574 (w/ load-balancing)
 - Wait until bucket _default17 is made
 - Run 4 ranged searches simultaneously spanning over the COLD buckets (half running on Node 8983, half running on Node 7574)
 - 7574 #1 -- latest: 2006-05-22T02:46:40Z, earliest: 2004-05-23T02:46:40Z
 - 8983 #1 -- latest: 2004-05-22T02:46:40Z, earliest: 2002-05-02T01:23:25Z
 - 8983 #2 -- latest: 2003-04-18T00:41:40Z, earliest: 2001-03-27T02:05:05Z
 - 7574 #2 -- latest: 2002-03-13T01:23:20Z, earliest: 2000-01-02T00:00:05Z
 - 8983 #3 – UNBOUNDED
 - 7574 #3 -- UNBOUNDED
 - Drop either Node 8983 or Node 7574 unexpectedly

- Re-run searches from dropped Node
- Expected outcome
 - Operation recovery
 - Dropped Node is Overseer
 - Failure of all operations
 - Re-election of New Overseer
 - New Overseer executes recovery from failed physical operations
 - Running Searches will re-hold buckets***(**not implemented**)
 - re-sending of operations to new overseer
 - Dropped Node is Slave
 - Completion or partial completion of physical operations utilizing slave node
 - Overseer Recovery from partially completed physical operations
 - Failure of all Search running on dropped Node
 - Overseer removes Bucket Holders originating from that Node
 - Failure of all Indexing on dropped Node
 - Replication will be paused, and Indexer replicas will no longer be indexed into
 - Continuation of Indexing and Searching on all live nodes
 - Indexing will continue, even though replication will not occur on the dead node
 - Failed Batches will be redirected to the live node
 - All data will be indexed
 - Searching will continue or fail, depending on whether a replica that it was searching on died, or whether a bucket cannot be physically held
 - New Searches on live nodes
 - Some will work, some will not because necessary buckets cannot be retrieved

(OPTIONAL) -- MORE DETAILED TESTS REGARDING OPERATION FAILURE & RECOVERY:

Operation Recovery from Overseer Death

- **CREATE BUCKET**
 - **TEST 1, 2, 3, & 4**
 - Procedure
 - Add breakpoint at IVrixBucket.createSelf(), on either line
 - createIsDetachedPropertyNode()
 - createBucketTypePropertyNode()
 - createBucket()
 - updateSolrMetadata()

- hit breakpoint, then fail overseer
- Expected Outcome
 - When new overseer boots, it will deem the bucket as dead/unusable, and will delete it
- **ROLL TO COLD**
 - **TEST 1 & 2**
 - Procedure
 - Add breakpoint at IVrixBucket.rollToCold(), on either line
 - prepareWarmBucketForColdRollover()
 - updateSolrMetadata()
 - hit breakpoint, then fail overseer
 - Expected Outcome
 - When new overseer boots, it will fully roll the bucket to COLD (i.e., remove all unnecessary replicas and update the solrMetadata field in the bucket)
- **ATTACH**
 - **TEST 1**
 - Procedure
 - Add breakpoint at IVrixBucket.physicallyAttachOnly()
 - hit breakpoint, then fail overseer
 - Expected Outcome
 - When new overseer boots, it will fully attach the bucket
- **DETACH**
 - **TEST 1**
 - Procedure
 - Add breakpoint at IVrixBucket.physicallyDetachOnly()
 - hit breakpoint, then fail overseer
 - Expected Outcome
 - When new overseer boots, it will fully detach the bucket

Operation Recovery from Slave Death

- **CREATE BUCKET**
 - **TEST 1**
 - Procedure
 - Add breakpoint at createAddReplicaRequest(...indexerNodeName...)
 - fail node, then release breakpoint
 - Expected Outcome
 - Bucket will fail to create, and will be a hanging/dead bucket
 - When overseer receives notification of dead node, it will find the hanging/dead bucket, and delete it
 - **TEST 2**
 - Procedure
 - Add breakpoint at createAddReplicaRequest(...replicationNodeName...)
 - fail node, then release breakpoint

- Expected Outcome
 - Replication replica will fail to be created
 - Bucket will continue to be created as normal, since failure to add replication replica does not impact indexing
- **ROLL TO COLD**
 - Cannot fail...
- **ATTACH**
 - **TEST 1**
 - Procedure
 - Add breakpoint at IVrixBucket.physicallyAttachOnly()
 - fail node, then release breakpoint
 - Expected Outcome
 - Failure to attach, but will not throw error
 - TRY TO HOLD BUCKET command will return “false”, since it could not get a physical hold of the bucket
- **DETACH**
 - Cannot fail...

7 – Recovery of Node after Node failure during Indexing + Searching on IVrix Cluster

- **TEST 1:**
 - Procedures:
 - {procedures from **section 6**}
 - Re-boot dropped Node
 - Run New Searches from that Node
 - Expected Outcome
 - Buckets
 - Replication Replicas will sync up with other replicas
 - Previous Hot Buckets should not exist
 - Indexer Replicas will have their state recovered (time bounds) *****(not implemented)**
 - State remains consistent, both on IVrixDB and Solr
 - currently executing Indexing and Searches will start to utilize recovered Node
 - New Searching commands can be executed that will utilize recovered Node
 - New Indexing commands can be executed that will utilize recovered Node